

Naïve Semantic Text Similarity Model

Zachary Parent

UPC

December 2, 2024

Outline

- 1 Introduction
- 2 Methodology
 - Approach
 - Feature Extraction
 - Model Training
 - Feature Selection
 - Model Evaluation
- 3 Results
- 4 Conclusions

Introduction

- Semantic Text Similarity (STS) is crucial for many NLP tasks
- Challenge: Which features best capture semantic similarity?
- Our approach: Unbiased feature analysis using Random Forests

Methodology

- Approach
- Feature extraction
- Feature selection
- Model training
- Model evaluation

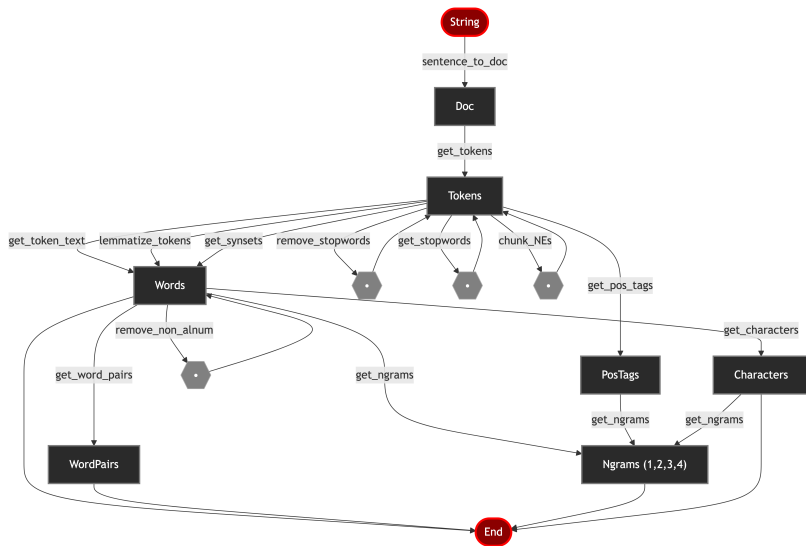
Approach

- Naïve approach which requires no knowledge of the corpus
- Use categorized steps to process sentences in every permutation
 - 520 permutations
 - e.g. `sentence_to_doc` → `chunk_NEs` → `remove_stopwords` → `lemmatize_tokens` → `get_characters` → `get_2grams`
- Apply 4 similarity metrics to each permutation
 - Jaccard, Cosine, Euclidean, Manhattan
- Used Random Forest's feature importance capabilities
- Let the data guide feature selection

Feature Extraction

Feature Extraction

Feature Extraction



generate_valid_permutations()

```
# Generate all valid permutations of sentence processing steps
def generate_valid_permutations(
    functions: List[Callable] = all_functions,
) -> List[Tuple[Callable, ...]]:
    valid_permutations = []
    for n in range(1, len(functions) + 1):
        for perm in itertools.permutations(functions, n):
            if _is_valid_permutation(perm):
                valid_permutations.append(perm)
    # Add sentence_to_doc to the beginning of each permutation
    valid_permutations = (
        tuple([sentence_to_doc] + perm for perm in valid_permutations))
    # Add final step to each permutation (e.g. get_2grams)
    valid_permutations = (
        [new_perm for perm in valid_permutations for new_perm in add_final_step(perm)])
    return valid_permutations
```



```
# Dictionary to hold function names and their input/output types
function_input_output_types: Dict[str, Tuple[Tuple[type, ...], type]] = {}

# Extract the input and output types of a function
def _extract_input_output_types(func: Callable) -> Tuple[type, type]:
    signature = inspect.signature(func)
    param_types = [param.annotation for param in signature.parameters.values()]
    return_type = signature.return_annotation
    return param_types[0], return_type

# Populate the dictionary with function names and their input/output types
for func in all_functions:
    input_types, output_type = _extract_input_output_types(func)
    function_input_output_types[func.__name__] = (input_types, output_type)

# Function to check if a permutation is valid based on input/output types
def _is_valid_permutation(perm: Tuple[Callable]) -> bool:
    if function_input_output_types[perm[0].__name__][0] != spacy.tokens.doc.Doc:
        return False
    if function_input_output_types[perm[-1].__name__][1] not in [
        Tuple[Word, ...],
        Tuple[PosTag, ...],
        Tuple[Character, ...],
    ]:
        return False
    for i in range(len(perm) - 1):
        _, current_func_output_type = function_input_output_types[perm[i].__name__]
        next_func_input_type, _ = function_input_output_types[perm[i + 1].__name__]
        if current_func_output_type != next_func_input_type:
            return False
    return True
```

```
class PosTag(str): pass
class Word(str): pass
class Character(str): pass
class Ngram(Tuple[Word | Character | PosTag, ...]): pass
class WordPair(Tuple[Word, Word]): pass

def get_characters(words: Tuple[Word, ...]) -> Tuple[Character, ...]:
def get_word_pairs(words: Tuple[Word, ...]) -> Tuple[WordPair, ...]:
def sentence_to_doc(sentence: str) -> spacy.tokens.doc.Doc:
def get_tokens(doc: spacy.tokens.doc.Doc) -> Tuple[spacy.tokens.token.Token, ...]:
def get_pos_tags(tokens: Tuple[spacy.tokens.token.Token, ...]) -> Tuple[PosTag, ...]:
def lemmatize_tokens(tokens: Tuple[spacy.tokens.token.Token, ...]) -> Tuple[Word, ...]:
def get_token_text(tokens: Tuple[spacy.tokens.token.Token, ...]) -> Tuple[Word, ...]:
def get_2grams(words: Tuple[Word | Character | PosTag, ...]) -> Tuple[Ngram, ...]:
def get_3grams(words: Tuple[Word | Character | PosTag, ...]) -> Tuple[Ngram, ...]:
def get_4grams(words: Tuple[Word | Character | PosTag, ...]) -> Tuple[Ngram, ...]:
def chunk_NEs(doc: spacy.tokens.doc.Doc) -> Tuple[spacy.tokens.token.Token, ...]:
def get_synsets(tokens: Tuple[spacy.tokens.token.Token, ...]) -> Tuple[Word, ...]:
def remove_non_alnum(words: Tuple[Word, ...]) -> Tuple[Word, ...]:
def remove_stopwords(tokens: Tuple[spacy.tokens.token.Token, ...])
-> Tuple[spacy.tokens.token.Token, ...]:
def get_stopwords(tokens: Tuple[spacy.tokens.token.Token, ...])
-> Tuple[spacy.tokens.token.Token, ...]:
```

Feature Extraction

```
lexical_functions = [  
    get_characters,      # Character-level patterns  
    get_tokens,         # Word tokenization  
    get_token_text,     # Raw word forms  
    remove_non_alnum,   # Character filtering  
    get_word_pairs,     # Word co-occurrences  
]  
  
semantic_functions = [  
    lemmatize_tokens,   # Normalize to base meaning  
    get_synsets,        # Word meanings/concepts  
    chunk_NEs,          # Named entity grouping  
    get_pos_tags,       # Part of speech (bridges lexical/semantic)  
]  
  
ngram_functions = [  
    get_2grams,         # Bigrams  
    get_3grams,         # Trigrams  
    get_4grams,         # 4-grams  
]  
  
preprocessing_functions = [  
    remove_stopwords,   # Filter non-content words  
    get_stopwords,      # Identify non-content words  
]  
  
all_functions = lexical_functions + semantic_functions + preprocessing_functions
```

Feature Extraction

Cache is king { \$, € }

```
@cache
def get_tokens(doc: spacy.tokens.doc.Doc) -> Tuple[spacy.tokens.token.Token, ...]:
    return tuple(token for token in doc)
```

- Many variations in ordering result in same output
- This is a prime candidate for dynamic programming
- Caching significantly speeds up feature extraction
- 520 processes x 4 metrics = 2080 features
 - across all ~5000 sentence pairs takes ~15 minutes

Model Training

- Random Forest
 - Built-in feature importance analysis
 - Bagging allows discovery of local patterns
 - Handles high-dimensional feature spaces well (2080 features)
 - Resistant to overfitting
- Model Configuration
 - 100 trees in ensemble
 - No max depth, since we are interested in identifying fine-grained feature importances

Feature Selection

- 520 permutations \times 4 metrics = 2080 features
- The random forest model allows us to inspect feature importances
- We used this to identify how important the top features are to the model
- We trained and validated models with different subsets of the top features

Model Evaluation

- Multi-level evaluation strategy
 - Initial 5-fold cross-validation on training set
 - 80/20 train/validation split for comparing feature sets
 - Held-out test set for final evaluation
- Custom Pearson correlation scorer
 - Measures linear correlation with gold standard
 - Implemented as sklearn-compatible scoring function
 - Allows direct comparison with published results

Results



Top Features

- Jaccard similarity dominates (7 of top 10)
- Common pipeline steps: lemmatization, remove stopwords, n-grams
- Top feature accounts for 20% importance

Feature Pipeline	Importance
score_jaccard_165	0.197
score_cosine_257	0.089
score_cosine_165	0.069
score_jaccard_258	0.033
score_cosine_258	0.022

Figure: Top 5 Features by Importance

Conclusions

- TODO